

Presented by

Dr. Trupti Padiya

**School of
Computing and
Creative
Technologies**

Advanced Databases UFCFU3-15-3

Introduction - NoSQL Databases

NoSQL databases

- NoSQL (Not Only SQL) or non-SQL databases refers to non relational databases. NoSQL databases store and manipulate data in other formats/data models than tabular relations/relational model. NoSQL databases do not require a fixed schema.
- Designed for horizontal scalability
- High availability
- Eventual consistency
- Support different data models. Some examples include key-value store, document store, graph databases etc.

Why NoSQL

- Massive data volumes
 - Massively distributed architecture required to store the data
 - Google, Amazon, Yahoo, Facebook – 10-100K servers
- Flexibility
 - Data model
 - Schema free design – easy to process structured and semi structured data e.g. JSON, XML
 - Schema evolution
- Performance
 - Extreme Query workload – complex query joins in RDBMS could be inefficient at scale

NoSQL Pros and Cons

■ Advantages

- Massive scalability
- High availability
- Lower cost (than competitive solutions at that scale)
- (usually) predictable elasticity
- Schema flexibility, sparse & semi-structured data

■ Disadvantages


- Limited query capabilities (so far)
- Eventual consistency is not intuitive to program for
- No standardization
- Portability might be an issue
- Insufficient access control

NoSQL Categorization

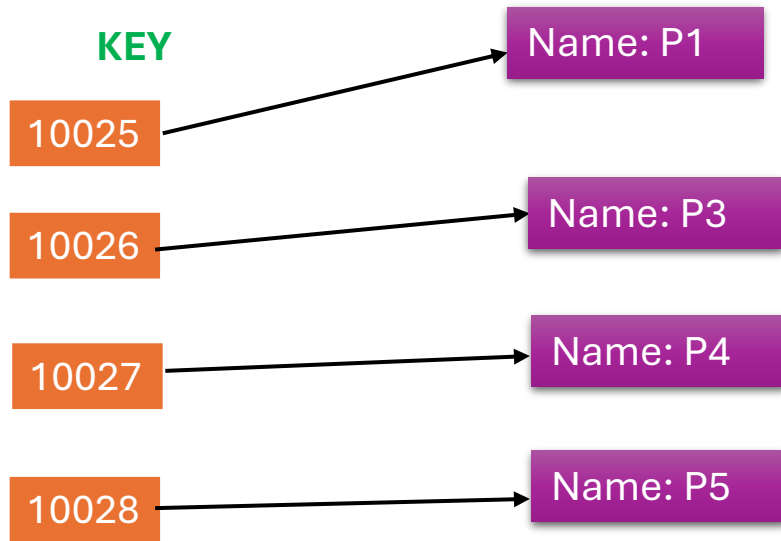
- Key-Value Stores
- Column Stores
- Document Stores
- Graph databases

Key-Value Stores

- Simple Data model –Key value pairs
- Key-value based database stores data as (key, value) pairs
 - Keys are unique
 - Hash map, hash table or dictionary



Key-Value Stores



- Keys are hashed by means of a so-called **hash function**
 - A hash function takes an arbitrary value of arbitrary size and maps it to a key with a fixed size, which is called the hash value.
 - Each hash can be mapped to a space in computer memory

Hash map/Hash tables

Key
10025
10026
10027
10028

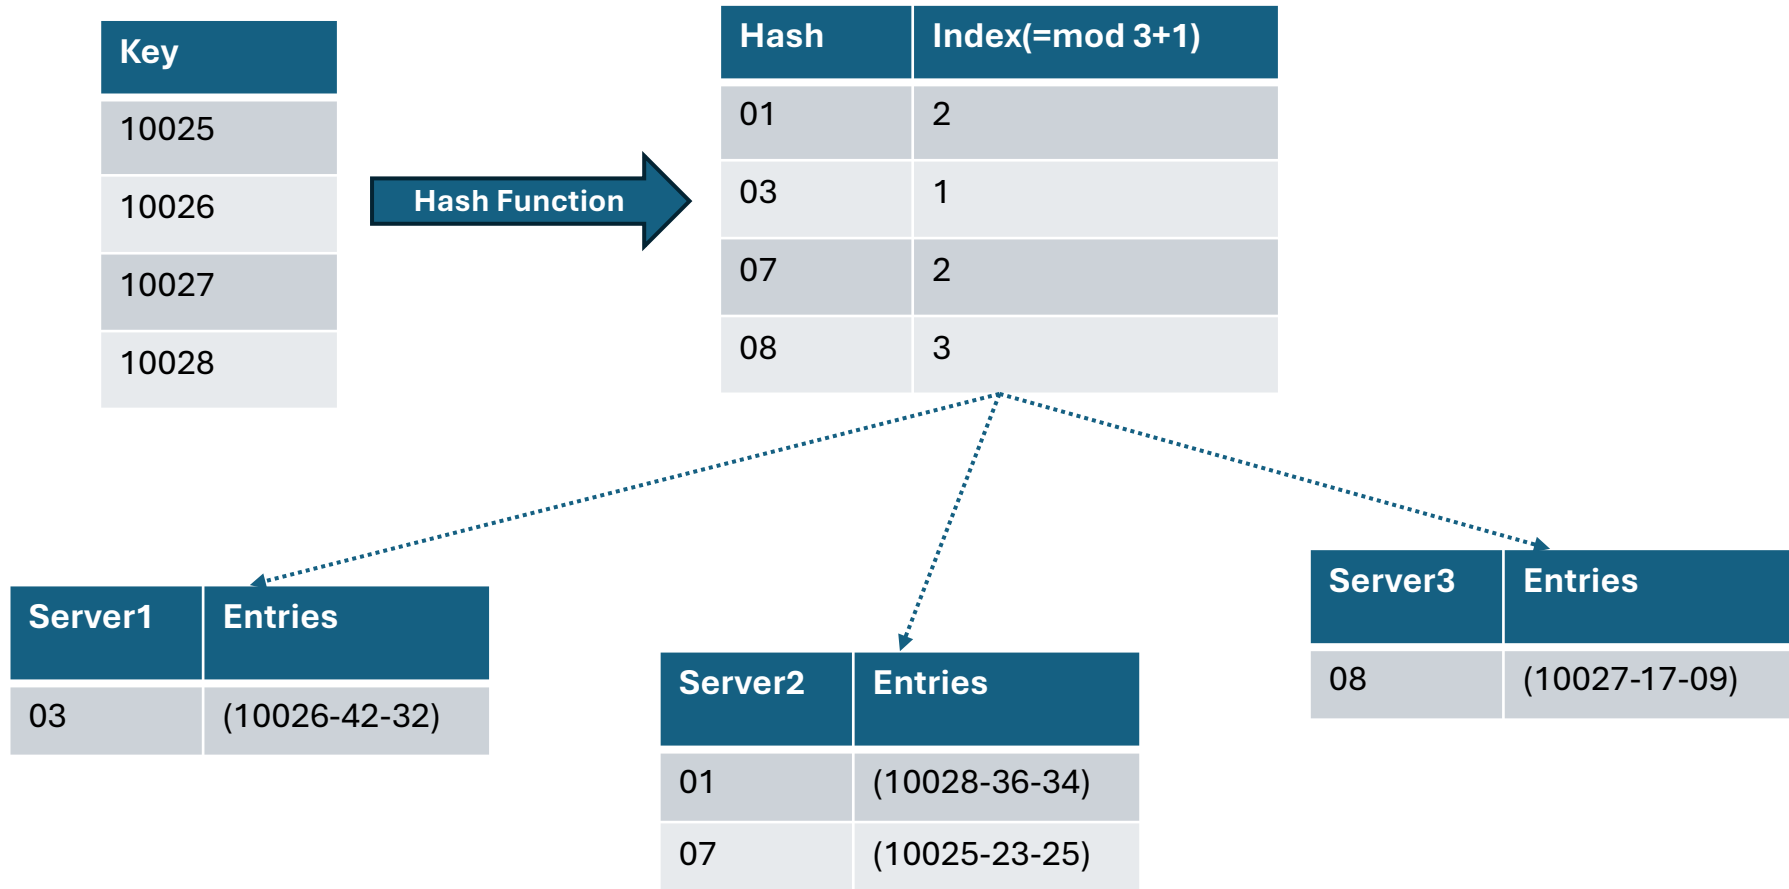


Hash	Key
01	(10028-36-34)
03	(10026-42-32)
07	(10025-23-25)
08	(10027-17-09)

Key-Value Store Sharding

- NoSQL databases are built with horizontal scalability support in mind
- Distribute hash table over different locations
- Assume we need to spread our hashes over three servers
 - Hash every key to a server identifier
 - $\text{index}(\text{hash}) = \text{mod}(\text{hash}, \text{nrServers}) + 1$
- Note: A database shard, or simply a shard, is a horizontal partition of data

Key-Value Store Sharding



$a \bmod b = a - (\text{Int}[a / b] \times b)$, where Int is an integer part of the value

Key-Value Stores Pros and Cons

■ Pros:

- simple model
- very fast
- very scalable
- able to distribute horizontally

■ Cons:

- many data structures (objects)
- cannot do complex queries
- key value pairs

Tuple Stores

- A **tuple store** is similar to a key-value store, with the difference that it does not store pairwise combinations of a key and a value, but instead stores a unique key together with a vector of data
- Example:
 - marc -> ("Marc", "McLast Name", 25, "Germany")
- Some Tuple Stores supports Schema-less design - No requirement to have the same length or semantic ordering
- Various NoSQL implementations do, however, permit organizing entries in semantical groups (aka collections or tables)

Tuple Stores Pros and Cons

■ Pros

- Efficient querying
- Support for Complex data structure
- Schema Flexibility
- Optimized for OLTP

■ Cons

- Performance issues with complex joins
- Generally not considered good for OLAP
- Write based operations

Column Stores

- A **column store** is a database management system that stores data tables as sections of columns of data
- Useful if
 - aggregates are regularly computed over large numbers of similar data items
 - data is sparse, i.e. columns with many null values

Column Stores

- Row based databases are not efficient at performing operations that apply to the entire data set
 - Need indexes which add overhead
- In a column-oriented database, all values of a column are placed together on disk

Genre: fantasy: 1, 4 education: 2, 3

Title: My first book:1 Beginners guide:2. SQL strikes back:3 The rise of SQL:4

Price: 20:1 10:2,4 40:3

Audiobook price: 30:1
- A column matches the structure of a normal index in a row-based system
- Operations such as: find all records with price equal to 10 can now be executed directly
- Null values do not take up storage space anymore

- **Example**

Id	Genre	Title	Price	Audiobook price
1	fantasy	My first book	20	30
2	education	Beginners guide	10	null
3	education	SQL strikes back	40	null
4	fantasy	The rise of SQL	10	null

Column Stores

■ Pros

- Schema Flexibility
- Analytical queries
- Data Compression
- Read based operations

■ Cons

- Write based operation
- Generally not considered good for OLTP
- Single row queries
- Join operations could be slowed down

Document Stores

- **Document stores** store a collection of attributes that are labeled and unordered, representing items that are semi-structured

- Example:

```
{  
  Title    = "Harry Potter"  
  ISBN     = "111-1111111111"  
  Authors  = ["J.K. Rowling"]  
  Price    = 32  
  Dimensions = "8.5 x 11.0 x 0.5"  
  PageCount = 234  
  Genre    = "Fantasy"  
}
```


Document Stores

- Most modern NoSQL databases choose to represent documents using JSON

```
{  
  "title": "Harry Potter",  
  "authors": ["J.K. Rowling", "R.J. Kowling"],  
  "price": 32.00,  
  "genres": ["fantasy"],  
  "dimensions": {  
    "width": 8.5,  
    "height": 11.0,  
    "depth": 0.5  
  },  
  "pages": 234,  
  "in_publication": true,  
  "subtitle": null  
}
```

Documents with keys

- Most NoSQL document stores will allow you to store documents in tables (collections) in a schema-less manner, but will enforce that a primary key be specified
 - e.g. Amazon's DynamoDB, MongoDB (`_id`)
- Primary key will be used as a partitioning key to create a hash and determine where the data will be stored

Document stores and complex queries

- Document stores do not support relations
- **First approach**: embedded documents

```
{
  "title": "Databases for Beginners",
  "authors": ["J.K. Sequel", "John Smith"],
  "pages": 234
}
{
  "title": "Databases for Beginners",
  "authors": [
    {"first_name": "Jay Kay", "last_name": "Sequel", "age": 54},
    {"first_name": "John", "last_name": "Smith", "age": 32}
  ],
  "pages": 234
}
```

But: Data duplication!

Document stores and complex queries

- **Second approach**: create two collections

book collection:

```
{  
  "title": "Databases for Beginners",  
  "authors": ["Jay Kay Sequel", "John Smith"],  
  "pages": 234  
}
```

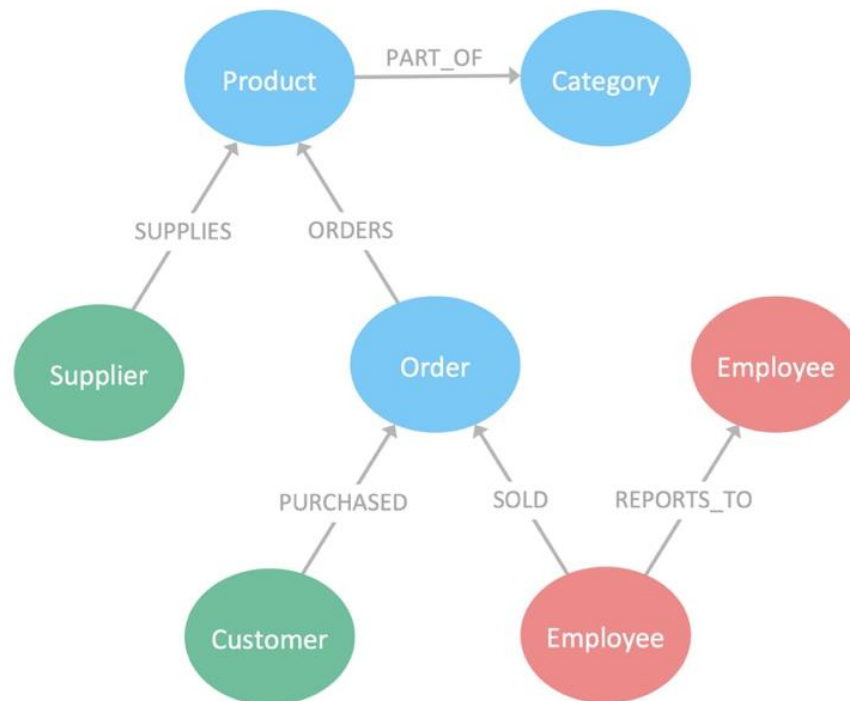
authors collection:

```
{  
  "_id": "Jay Kay Sequel",  
  "age": 54  
}
```

But: Need to resolve complex relational queries in application code!

Graph databases

- **Graph databases** apply graph theory to the storage and retrieval of data
- Graphs consist of **nodes** and **edges**



Some (extra) resources/links:

- <https://www.mongodb.com/resources/basics/databases/document-databases>
- <https://www.couchbase.com/blog/columnar-store-vs-row-store/>
- <https://neo4j.com/docs/getting-started/graph-database/>
- <https://www.mongodb.com/resources/basics/databases/key-value-database#:~:text=Key%20value%20databases%2C%20also%20known,associated%20value%20with%20each%20key.>